

# Business Case: LoanTap Logistic Regression

## Problem Statement

LoanTap is an online platform committed to delivering customized loan products to millennials. They innovate in an otherwise dull loan segment, to deliver instant, flexible loans on consumer friendly terms to salaried professionals and businessmen.

The data science team at LoanTap is building an underwriting layer to determine the creditworthiness of MSMEs as well as individuals.

LoanTap deploys formal credit to salaried individuals and businesses 4 main financial instruments:

Personal Loan EMI Free Loan Personal Overdraft Advance Salary Loan This case study will focus on the underwriting process behind Personal Loan only

### Problem Statement:

Given a set of attributes for an Individual, determine if a credit line should be extended to them. If so, what should the repayment terms be in business recommendations.

Predicts the who can pay loan fully and who can not.

```
In [697]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
import re
from sklearn.model_selection import train_test_split, KFold, cross_val_score
from sklearn.preprocessing import StandardScaler
from imblearn.over_sampling import SMOTE
from sklearn.linear_model import LogisticRegression
from sklearn.pipeline import make_pipeline
from sklearn.metrics import roc_curve, roc_auc_score
from sklearn.metrics import precision_recall_curve
from sklearn.metrics import auc
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.model_selection import cross_val_score
```

```
In [613]: data = pd.read_csv('loan.csv')
```

In [614]: `data.head()`

Out[614]:

	loan_amnt	term	int_rate	installment	grade	sub_grade	emp_title	emp_length	home
0	10000.0	36 months	11.44	329.48	B	B4	Marketing	10+ years	
1	8000.0	36 months	11.99	265.68	B	B5	Credit analyst	4 years	
2	15600.0	36 months	10.49	506.97	B	B3	Statistician	< 1 year	
3	7200.0	36 months	6.49	220.65	A	A2	Client Advocate	6 years	
4	24375.0	60 months	17.27	609.33	C	C5	Destiny Management Inc.	9 years	

5 rows × 27 columns

## Exploratory Data Analysis

In [615]: `data.shape`

Out[615]: (396030, 27)

In [616]: data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 396030 entries, 0 to 396029
Data columns (total 27 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   loan_amnt                             396030 non-null float64
1   term                                  396030 non-null object
2   int_rate                              396030 non-null float64
3   installment                           396030 non-null float64
4   grade                                 396030 non-null object
5   sub_grade                             396030 non-null object
6   emp_title                             373103 non-null object
7   emp_length                            377729 non-null object
8   home_ownership                        396030 non-null object
9   annual_inc                            396030 non-null float64
10  verification_status                   396030 non-null object
11  issue_d                               396030 non-null object
12  loan_status                           396030 non-null object
13  purpose                               396030 non-null object
14  title                                 394275 non-null object
15  dti                                    396030 non-null float64
16  earliest_cr_line                      396030 non-null object
17  open_acc                              396030 non-null float64
18  pub_rec                               396030 non-null float64
19  revol_bal                             396030 non-null float64
20  revol_util                            395754 non-null float64
21  total_acc                             396030 non-null float64
22  initial_list_status                   396030 non-null object
23  application_type                      396030 non-null object
24  mort_acc                              358235 non-null float64
25  pub_rec_bankruptcies                  395495 non-null float64
26  address                               396030 non-null object
dtypes: float64(12), object(15)
memory usage: 81.6+ MB
```

In [617]: data.describe()

Out[617]:

	loan_amnt	int_rate	installment	annual_inc	dti	open
count	396030.000000	396030.000000	396030.000000	3.960300e+05	396030.000000	396030.00
mean	14113.888089	13.639400	431.849698	7.420318e+04	17.379514	11.31
std	8357.441341	4.472157	250.727790	6.163762e+04	18.019092	5.13
min	500.000000	5.320000	16.080000	0.000000e+00	0.000000	0.00
25%	8000.000000	10.490000	250.330000	4.500000e+04	11.280000	8.00
50%	12000.000000	13.330000	375.430000	6.400000e+04	16.910000	10.00
75%	20000.000000	16.490000	567.300000	9.000000e+04	22.980000	14.00
max	40000.000000	30.990000	1533.810000	8.706582e+06	9999.000000	90.00

```
In [618]: print("Percentage of missing value")
np.round((data.isna().sum()/data.shape[0])*100,2)
```

Percentage of missing value

```
Out[618]: loan_amnt      0.00
term      0.00
int_rate  0.00
installment 0.00
grade     0.00
sub_grade 0.00
emp_title  5.79
emp_length 4.62
home_ownership 0.00
annual_inc 0.00
verification_status 0.00
issue_d    0.00
loan_status 0.00
purpose    0.00
title      0.44
dti        0.00
earliest_cr_line 0.00
open_acc   0.00
pub_rec    0.00
revol_bal  0.00
revol_util 0.07
total_acc  0.00
initial_list_status 0.00
application_type 0.00
mort_acc   9.54
pub_rec_bankruptcies 0.14
address    0.00
dtype: float64
```

## • Univariate Analysis

```
In [619]: for i in data.columns:
          print(i,":",data[i].nunique())
```

```
loan_amnt : 1397
term : 2
int_rate : 566
installment : 55706
grade : 7
sub_grade : 35
emp_title : 173105
emp_length : 11
home_ownership : 6
annual_inc : 27197
verification_status : 3
issue_d : 115
loan_status : 2
purpose : 14
title : 48817
dti : 4262
earliest_cr_line : 684
open_acc : 61
pub_rec : 20
revol_bal : 55622
revol_util : 1226
total_acc : 118
initial_list_status : 2
application_type : 3
mort_acc : 33
pub_rec_bankruptcies : 9
address : 393700
```

```
In [620]: for i in data.columns:
          print(i)
          print(data[i].value_counts())
          print("_____")
```

```
loan_amnt
10000.0    27668
12000.0    21366
15000.0    19903
20000.0    18969
35000.0    14576
...
36625.0      1
37450.0      1
36275.0      1
38225.0      1
725.0        1
Name: loan_amnt, Length: 1397, dtype: int64
```

---

```
term
36 months    302005
60 months     94025
Name: term, dtype: int64
```

---

```
In [621]: np.round((data['loan_status'].value_counts()/data.shape[0])*100,2)
```

```
Out[621]: Fully Paid      80.39  
Charged Off    19.61  
Name: loan_status, dtype: float64
```

```
In [622]: np.round((data['home_ownership'].value_counts()/data.shape[0])*100,2)
```

```
Out[622]: MORTGAGE      50.08  
RENT        40.35  
OWN         9.53  
OTHER       0.03  
NONE        0.01  
ANY         0.00  
Name: home_ownership, dtype: float64
```

```
In [623]: np.round((data[(data.grade=='A') & (data.loan_status=="Fully Paid")].shape[0]/data.shape[0])*100,2)
```

```
Out[623]: 93.71
```

```
In [624]: data_cat = []  
data_num = []  
for i in data.columns:  
    if data[i].nunique()<50:  
        data_cat.append(i)  
    else:  
        data_num.append(i)
```

```
In [625]: data_cat
```


```
Out[625]: ['term',  
'grade',  
'sub_grade',  
'emp_length',  
'home_ownership',  
'verification_status',  
'loan_status',  
'purpose',  
'pub_rec',  
'initial_list_status',  
'application_type',  
'mort_acc',  
'pub_rec_bankruptcies']
```

```
In [626]: data_num.remove('issue_d')  
data_num.remove('emp_title')  
data_num.remove('earliest_cr_line')  
data_num.remove('address')  
data_num.remove('title')
```


```
In [627]: data_num
```

```
Out[627]: ['loan_amnt',  
           'int_rate',  
           'installment',  
           'annual_inc',  
           'dti',  
           'open_acc',  
           'revol_bal',  
           'revol_util',  
           'total_acc']
```


```
In [628]: for i in data_cat:  
           plt.figure(i,figsize=(15,8))  
           sns.countplot(x=i,data=data)
```



```
In [629]: for i in data_num:
plt.figure(i,figsize=(15,8))
sns.distplot(data[i])
```

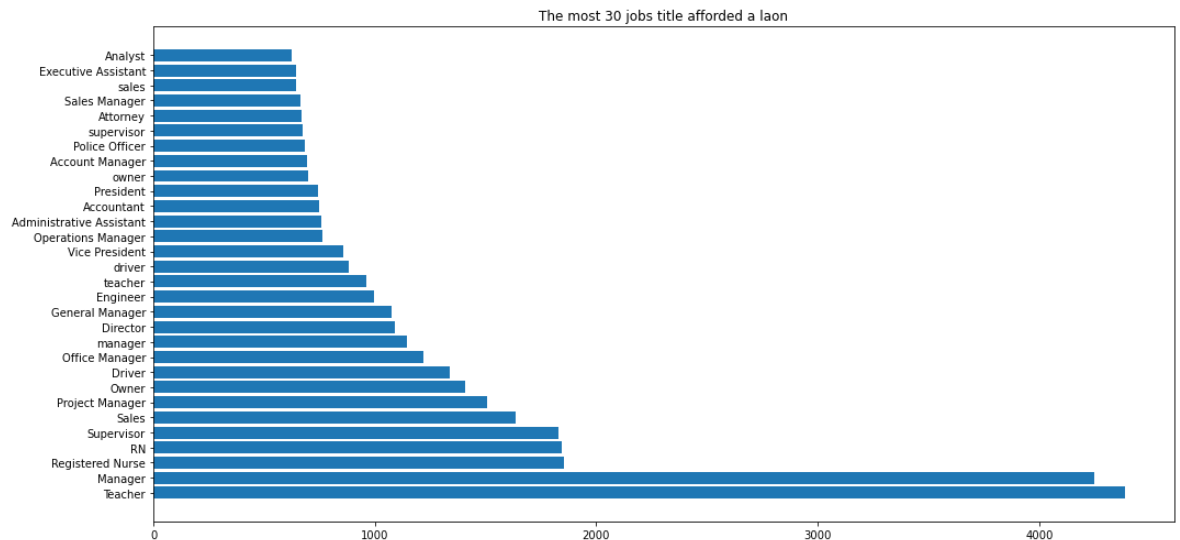


```
In [630]: for i in data_num:
plt.figure(i,figsize=(15,8))
sns.boxplot(data=data,x=i)
```





```
In [631]: plt.figure(figsize=(15,7))
plt.barh(data.emp_title.value_counts()[ :30].index,data.emp_title.value_coun
plt.title("The most 30 jobs title afforded a laon")
plt.tight_layout()
```

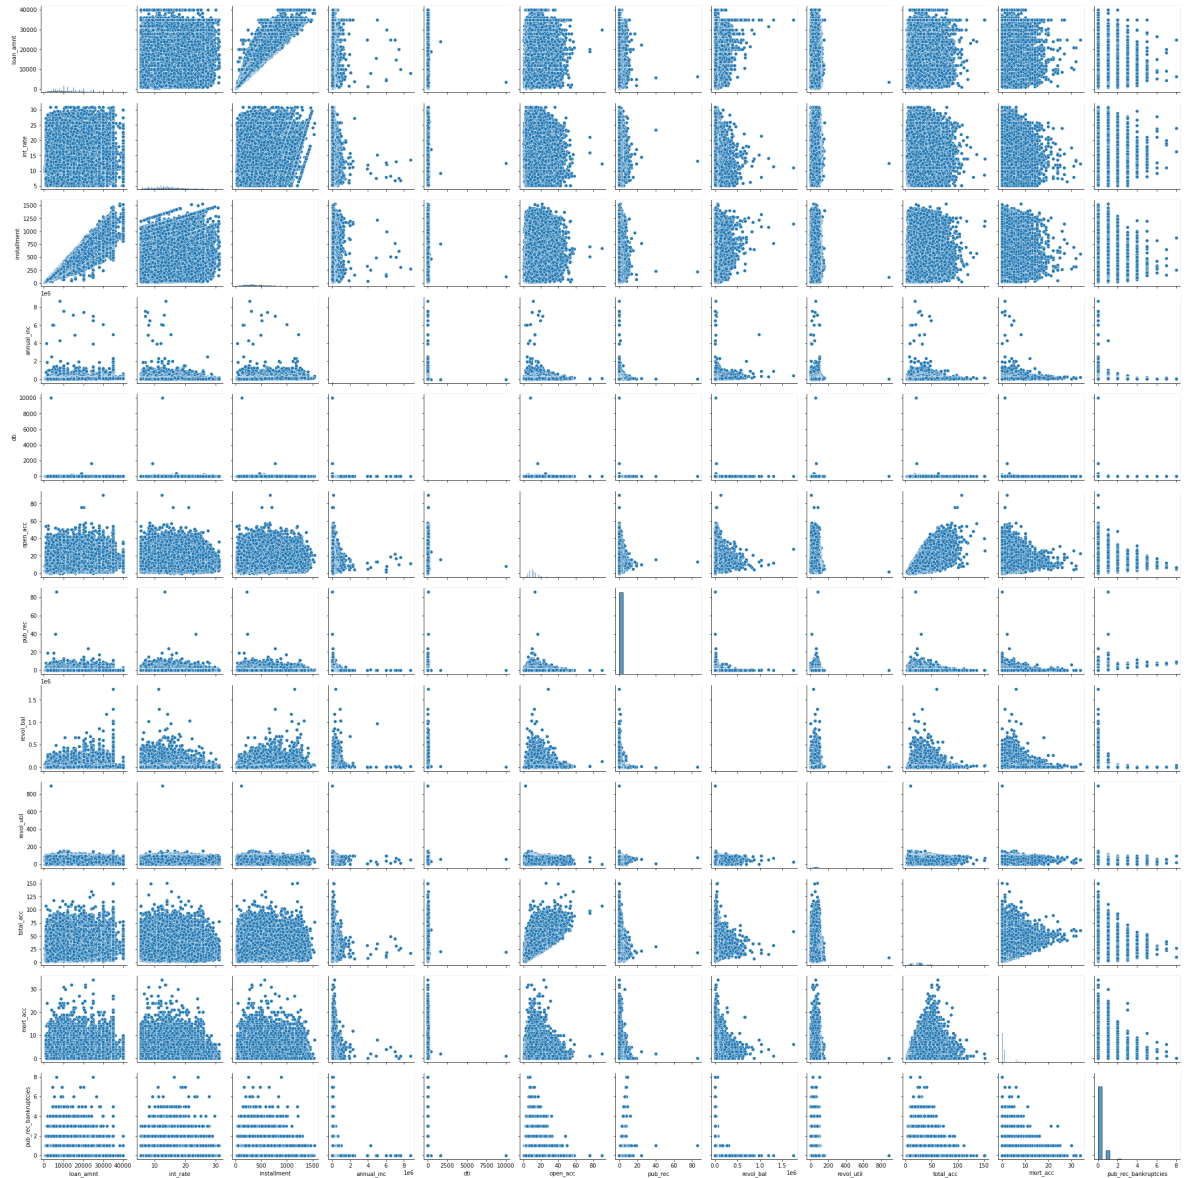


- dataset is highly imbalanced.
- int\_rate, installments, annual\_inc, dti, revol\_acc, total\_acc, open\_acc have very high number outliers.
- Maximum people taking 36 month plan
- 10+ years of experience person taking more loan
- Mortgage and Rent home ownership person have dominant number
- All numerical feature are right skewed.

## • Bivariate Analysis


```
In [632]: sns.pairplot(data)
```

```
Out[632]: <seaborn.axisgrid.PairGrid at 0x296ae064c10>
```




```
In [633]: sns.scatterplot(data['loan_amnt'],data['installment'])
```

```
Out[633]: <AxesSubplot:xlabel='loan_amnt', ylabel='installment'>
```




```
In [634]: for i in data_num:
plt.figure(i,figsize=(15,8))
sns.boxplot(x=i,y='loan_status',data=data,palette='rainbow')
```



```
In [635]: plt.figure(figsize=(15,8))
sns.heatmap(data.corr(), cmap="YlGnBu", annot=True)
```

Out[635]: <AxesSubplot:>



- Installment and loan amount are highly correlated to each other.
- 13% rate of interest is 50 percentile for fully paid and 25 percentile for charged off.

## Data Preprocessing

```
In [636]: data=data.drop(["emp_title","issue_d","title","earliest_cr_line","sub_grade"])
```

```
In [637]: data=data.drop(['installment'],axis=1)
```

## 1. Duplicate value check

```
In [638]: data[data.duplicated()].sum()
```

```
Out[638]: loan_amnt      0.0
term      0.0
int_rate  0.0
grade     0.0
home_ownership  0.0
annual_inc  0.0
verification_status  0.0
loan_status  0.0
purpose   0.0
dti       0.0
open_acc  0.0
pub_rec   0.0
revol_bal  0.0
revol_util  0.0
total_acc  0.0
initial_list_status  0.0
application_type  0.0
mort_acc   0.0
pub_rec_bankruptcies  0.0
address    0.0
dtype: float64
```

No duplicates value is present

## 2. Missing value treatment

```
In [639]: data.isna().sum()
```

```
Out[639]: loan_amnt      0
term      0
int_rate  0
grade     0
home_ownership  0
annual_inc  0
verification_status  0
loan_status  0
purpose   0
dti       0
open_acc  0
pub_rec   0
revol_bal  0
revol_util 276
total_acc  0
initial_list_status  0
application_type  0
mort_acc   37795
pub_rec_bankruptcies  535
address    0
dtype: int64
```

```
In [640]: median = data['revol_util'].median()
data['revol_util'] = data['revol_util'].fillna(median)
```

filling null values of revol\_util with the median value

```
In [641]: data['mort_acc'] = data['mort_acc'].fillna(1)
data['pub_rec_bankruptcies'] = data['pub_rec_bankruptcies'].fillna(1)
```

filling null value of mort\_acc and pub\_rec\_bankruptcies with the "1" because "0" is greater in size

```
In [642]: data.isna().sum()
```

```
Out[642]: loan_amnt      0
term      0
int_rate  0
grade     0
home_ownership  0
annual_inc  0
verification_status  0
loan_status  0
purpose   0
dti       0
open_acc  0
pub_rec   0
revol_bal  0
revol_util  0
total_acc  0
initial_list_status  0
application_type  0
mort_acc   0
pub_rec_bankruptcies  0
address    0
dtype: int64
```

### 3. Outlier treatment

```
In [689]: numerical_data=data.select_dtypes(include='number')
num_cols=numerical_data.columns
```

```
In [644]: for col in num_cols:
    mean=data[col].mean()
    std=data[col].std()

    upper_limit=mean+3*std
    lower_limit=mean-3*std

    data=data[(data[col]<upper_limit) & (data[col]>lower_limit)]

data.shape
```

```
Out[644]: (365525, 20)
```

taking three time of standard deviation more and less from mean as normal data apart from that removing because they are outliers.

## 4. Feature engineering

```
In [645]: data['pub_rec'] = data['pub_rec'].apply(lambda x: 1 if x >= 1 else 0)
data['mort_acc'] = data['mort_acc'].apply(lambda x: 1 if x >= 1 else 0)
data['pub_rec_bankruptcies'] = data['pub_rec_bankruptcies'].apply(lambda x:
```

```
In [646]: # Mapping of target variable
data['loan_status'] = data.loan_status.map({'Fully Paid':0, 'Charged Off':1})
```

```
In [647]: term_values = {' 36 months': 36, ' 60 months': 60}
data['term'] = data.term.map(term_values)
```

```
In [648]: list_status = {'w': 0, 'f': 1}
data['initial_list_status'] = data.initial_list_status.map(list_status)
```

```
In [649]: data['zip_code'] = data.address.apply(lambda x: x[-5:])
```

```
In [650]: data['zip_code'].value_counts()
```

```
Out[650]: 70466      52561
30723      52232
22690      52131
48052      51681
00813      42210
29597      41904
05113      41878
11650      10384
93700      10335
86630      10209
Name: zip_code, dtype: int64
```

```
In [651]: data.columns
```

```
Out[651]: Index(['loan_amnt', 'term', 'int_rate', 'grade', 'home_ownership',
               'annual_inc', 'verification_status', 'loan_status', 'purpose', 'dti',
               'open_acc', 'pub_rec', 'revol_bal', 'revol_util', 'total_acc',
               'initial_list_status', 'application_type', 'mort_acc',
               'pub_rec_bankruptcies', 'address', 'zip_code'],
              dtype='object')
```

```
In [652]: dummies_col = ["grade", "home_ownership", "zip_code", "verification_status", "dti"]
data = pd.get_dummies(data, columns=dummies_col, drop_first=True)
```

one hot encoding


```
In [653]: data=data.drop(['address'],axis=1)
```

```
In [654]: data.head()
```

```
Out[654]:
```

	loan_amnt	term	int_rate	annual_inc	loan_status	dti	open_acc	pub_rec	revol_bal	rev
0	10000.0	36	11.44	117000.0	0	26.24	16.0	0	36369.0	
1	8000.0	36	11.99	65000.0	0	22.05	17.0	0	20131.0	
2	15600.0	36	10.49	43057.0	0	12.79	13.0	0	11987.0	
3	7200.0	36	6.49	54000.0	0	2.60	6.0	0	5472.0	
4	24375.0	60	17.27	55000.0	1	33.95	13.0	0	24584.0	

5 rows × 51 columns



```
In [655]: data.shape
```

```
Out[655]: (365525, 51)
```

## 5. Data preparation for modeling

```
In [656]: X = data.drop('loan_status',axis=1)
          y = data[['loan_status']]
```

```
In [657]: X_tr_cv, X_test, y_tr_cv, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
          X_train, X_val, y_train, y_val = train_test_split(X_tr_cv, y_tr_cv, test_size=0.2, random_state=42)
```

```
In [658]: print(X_train.shape,X_test.shape,X_val.shape)
          print(y_train.shape,y_test.shape,y_val.shape)
```

```
(219315, 50) (73105, 50) (73105, 50)
(219315, 1) (73105, 1) (73105, 1)
```

```
In [659]: sm=SMOTE(random_state=42)
          X_train,y_train=sm.fit_resample(X_train,y_train)
```

Using Smote technique for oversampling the data.

```
In [660]: print('After OverSampling, the shape of train_X: {}'.format(X_train.shape))
          print('After OverSampling, the shape of train_y: {}'.format(y_train.shape))

          print("After OverSampling, counts of label 1: {}".format(sum(y_train['loan_status']==1)))
          print('After OverSampling, counts of label 0: {}'.format(sum(y_train['loan_status']==0)))
```

```
After OverSampling, the shape of train_X: (352342, 50)
After OverSampling, the shape of train_y: (352342, 1)
```

```
After OverSampling, counts of label 1: 176171
After OverSampling, counts of label 0: 176171
```



```
In [661]: sc = StandardScaler()
sc.fit(X_train)
X_train = sc.transform(X_train)
X_test = sc.transform(X_test)
X_val = sc.transform(X_val)
```


```
In [662]: def accuracy(y_true, y_pred):
y_true = y_true.values.reshape(len(y_true))
return np.sum(y_true==y_pred)/y_true.shape[0]
```

```
In [663]: train_scores = []
val_scores = []
sc = StandardScaler()
for la in np.arange(0.01, 5000.0, 100): # range of values of Lambda
scaled_lr = make_pipeline(sc, LogisticRegression(C=1/la))
scaled_lr = make_pipeline(sc, LogisticRegression(C=1/la))
scaled_lr.fit(X_train, y_train)
train_score = accuracy(y_train, scaled_lr.predict(X_train))
val_score = accuracy(y_val, scaled_lr.predict(X_val))
train_scores.append(train_score)
val_scores.append(val_score)
```

Checking the accurate lambda value for the model with the help of validation data

```
In [664]: plt.figure(figsize=(10,5))
plt.plot(list(np.arange(0.01, 5000.0, 100)), train_scores, label="train")
plt.plot(list(np.arange(0.01, 5000.0, 100)), val_scores, label="val")
plt.legend(loc='lower right')

plt.xlabel("Regularization Parameter( $\lambda$ )")
plt.ylabel("Accuracy")
plt.grid()
plt.show()
```



## Model building

```
In [665]: model=LogisticRegression(C=1/500)
          model.fit(X_train,y_train)
```

Out[665]: LogisticRegression(C=0.002)

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [666]: print('Accuracy of Logistic Regression Classifier on train set: {:.3f}'.for
          print('Accuracy of Logistic Regression Classifier on validation set: {:.3f]
          print('Accuracy of Logistic Regression Classifier on test set: {:.3f}'.forr
```

```
Accuracy of Logistic Regression Classifier on train set: 0.907
Accuracy of Logistic Regression Classifier on validation set: 0.868
Accuracy of Logistic Regression Classifier on test set: 0.867
```

```
In [667]: model.coef_
```

Out[667]: array([[ 0.15447456, 0.24230291, 1.26007837, -0.15551113, 0.20939823,  
 0.10091532, 0.02072214, -0.07833206, 0.15290854, -0.05577034,  
 -0.19279223, -0.18290777, -0.07440586, -0.67847981, -0.8279282 ,  
 -0.89017751, -0.7963395 , -0.5959547 , -0.28532041, -1.50132455,  
 -0.02388497, -0.04838529, -0.80181447, -1.43762894, -1.14529903,  
 0.87522784, -0.17308467, -1.14963745, -0.17499601, -0.15694041,  
 -0.16992693, 0.85844535, 0.87782022, -0.18353282, -0.24633107,  
 -0.96019982, -1.18385376, -0.04226879, -0.4871681 , -0.16501678,  
 -0.31910162, -0.2114169 , -0.17613155, -0.47276147, -0.05642384,  
 -0.19897519, -0.17062446, -0.17851618, -0.01629657, -0.0485730  
 4]))

```
In [668]: model.intercept_
```


Out[668]: array([0.73371401])

```
In [707]: # checking cross validation score
          score_lr=cross_val_score(LogisticRegression(), X_val, y_val,cv=5)
          print(score_lr)
          print("Avg :",np.average(score_lr))
```

```
[0.89159428 0.8861227  0.88892689 0.88974762 0.88858491]
Avg : 0.8889952807605498
```

```
In [674]: coefficients = model.coef_
avg_importance = np.mean(np.abs(coefficients), axis=0)
feature_importance = pd.DataFrame({'Feature': X, 'Importance': avg_importance})
feature_importance = feature_importance.sort_values('Importance', ascending=False)
feature_importance.plot(x='Feature', y='Importance', kind='barh', figsize=(10, 5))
```

```
Out[674]: <AxesSubplot:ylabel='Feature'>
```



# Results Evaluation

## 1. ROC AUC Curve & comments

```
In [676]: probability = model.predict_proba(X_test)
```


```
In [677]: probability
```

```
Out[677]: array([[0.29346712, 0.70653288],
 [0.00277887, 0.99722113],
 [0.99761274, 0.00238726],
 ...,
 [0.99795889, 0.00204111],
 [0.69728686, 0.30271314],
 [0.90768234, 0.09231766]])
```

```
In [678]: probabilites = probability[:,1]
```

```
In [679]: fpr, tpr, thr = roc_curve(y_test,probabilites)
```

```
In [680]: plt.plot(fpr,tpr)
plt.plot(fpr,fpr,'--',color='red' )
plt.title('ROC curve')
plt.xlabel('FPR')
plt.ylabel('TPR')
plt.show()
```



```
In [681]: # AUC
roc_auc_score(y_test,probabilites)
```

```
Out[681]: 0.8416642860615695
```

ROC AUC score is pretty good that means our model is predicting 1 to 1 and 0 to 0 correctly.


An ROC curve plots TPR vs. FPR at different classification thresholds. Lowering the classification threshold classifies more items as positive, thus increasing both False Positives

## 2. Precision Recall Curve & comments

```
In [682]: precision, recall, thr = precision_recall_curve(y_test, probabilites)
```

```
In [683]: plt.plot(recall, precision)

plt.xlabel('Recall')
plt.ylabel('Precision')
plt.title('PR curve')
plt.show()
```




```
In [684]: def precision_recall_curve_plot(y_test, pred_proba_c1):
    precisions, recalls, thresholds = precision_recall_curve(y_test, pred_proba_c1)

    threshold_boundary = thresholds.shape[0]
    #plot precision
    plt.plot(thresholds, precisions[0:threshold_boundary], linestyle='--', label='precision')
    #plot recall
    plt.plot(thresholds, recalls[0:threshold_boundary], label='recalls')

    start, end = plt.xlim()
    plt.xticks(np.round(np.arange(start, end, 0.1), 2))

    plt.xlabel('Threshold Value')
    plt.ylabel('Precision and Recall Value')
    plt.legend()
    plt.grid()
    plt.show()

precision_recall_curve_plot(y_test, logreg.predict_proba(X_test)[: , 1])
```



Precision-Recall is a useful measure of success of prediction when the classes are very imbalanced. In information retrieval, precision is a measure of result relevancy, while recall is a measure of how many truly relevant results are returned.

The precision-recall curve shows the tradeoff between precision and recall for different threshold. A high area under the curve represents both high recall and high precision, where high precision relates to a low false positive rate, and high recall relates to a low false negative rate. High scores for both show that the classifier is returning accurate results (high precision), as well as returning a majority of all positive results (high recall).

Curve is looking good.


### 3. Classification Report (Confusion Matrix etc)

```
In [693]: y_pred = model.predict(X_train)

          conf_matrix = confusion_matrix(y_train, y_pred)

          ConfusionMatrixDisplay(conf_matrix).plot()
```

Out[693]: <sklearn.metrics.\_plot.confusion\_matrix.ConfusionMatrixDisplay at 0x29751c055e0>




```
In [691]: y_pred = model.predict(X_test)

          conf_matrix = confusion_matrix(y_test, y_pred)

          ConfusionMatrixDisplay(conf_matrix).plot()
```

Out[691]: <sklearn.metrics.\_plot.confusion\_matrix.ConfusionMatrixDisplay at 0x29751ba13a0>




```
In [692]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.88	0.96	0.92	58723
1	0.76	0.48	0.59	14382
accuracy			0.87	73105
macro avg	0.82	0.72	0.75	73105
weighted avg	0.86	0.87	0.86	73105

```
In [695]: Precision =(((conf_matrix.T)/(conf_matrix.sum(axis=1))).T)
print("Precision Score: ",precision_score(y_train, y_pred))
print("Precision matrix")
plt.figure(figsize=(10,7))
sns.heatmap(Precision, annot=True, cmap="YlGnBu", fmt=".3f")
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()
```

Precision Score: 0.9599581681230832

Precision matrix






```
In [696]: print("Recall Score: ", recall_score(y_train, y_pred))
Recall = (conf_matrix / conf_matrix.sum(axis=0))
print("Recall matrix")
plt.figure(figsize=(10, 7))
sns.heatmap(Recall, annot=True, cmap="YlGnBu", fmt=".3f")
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()
```

Recall Score: 0.8492941517048777

Recall matrix



#### 4. Tradeoff Questions:

**1. How can we make sure that our model can detect real defaulters and there are less false positives? This is important as we can lose out on an opportunity to finance more individuals and earn interest on it.**

=> To keep very less False Positives, oversampling techniques like SMOTE should be used in model creation. Also we can use more complex algorithms like SVM, Decision-Trees, Random Forest and also try various hyperparameter tuning.

=> As you can see from the data, the percentage of defaulters is slightly higher than Banking industry.

**2. Since NPA (non-performing asset) is a real problem in this industry, it's important we play safe and shouldn't disburse loans to anyone.**

=> Yes. LoanTap should not disburse loans to everyone. Company's internal policy and analysis should be in place to identify the correct persons. From data provided, 20% of people default on their loan, which become NPAs for the company.

=> Low False positive means we should create the model with high Precision values. This can be achieved if we are keeping high threshold value in logistic Regression model.

=> But keeping too high values for threshold will increase False Negatives. This intuition may result in opportunity loss. In this case we will not give loans to persons which will not default but our model has predicted that they will default.

## Actionable Insights & Recommendations

- Around 80.39% of customers have fully paid their Loan Amount. The defaulters are ~ 20%. From Personal loan business perspective this ratio is high. These 20% will contribute in NPAs of LoanTap. To reduce the risk of NPAs,
- LoanTap should add slightly stringent rules to bring down this ratio to 5% to 6%.
- LoanTap should provide loans at slightly higher rate than other Banks. This will offset the risks of defaulters and maintain the profitability of the business.

- **Overall Statistics of the Model:**

- Accuracy = 87%
- Precision = 88
- Recall = 96
- F1 -score = 92%

- Model created has high values for accuracy, precision, recall & f1-score. This means, this model is a good classifier. Overall, it has good prediction capability in identifying right customers (which can be easily converted).
- Those person who has high rate of interest, they charged off. Company should give some relaxation in rate of interest.
- However this model has slightly low capability on correctly identifying defaulters. Overall data has 20% defaulters, model is able to predict 10% of them correctly.
- Using this model, LoanTap can easily reduce the ration of defaulters in their portfolio.
- application\_type JOINT has positive Coefficient. Which means LoanTap can promote persons to apply for joint loan. Because of this, chances of default will reduce.
- Purpose Renewable energy has negative Coefficient. This means LoanTap should stick to giving loans to conventional purposes like Marriage, car etc.
- term 60 months has negative Coefficient. Which means LoanTap should focus more on Loans for shorter duration (i.e. 36 months). Their social media campaign and marketing strategy should be based on this consideration.

## Questionnaire

### 1. What percentage of customers have fully paid their Loan Amount?

80.39%

### 2. Comment about the correlation between Loan Amount and Installment

the spearman correlation coefficient between Loan Amount and Installment features is very high. It is around 0.95.

**3. The majority of people have home ownership as \_\_\_\_.**

Majority of people have home ownership as Mortgage and Rent.

- MORTGAGE 198348 (~50.08%)
- RENT 159790 (~40.35)

**4. People with grades 'A' are more likely to fully pay their loan. (T/F)**

Yes, Out of all people with grade 'A', 93.7% paid loan fully.

**5. Name the top 2 afforded job titles.**

Teacher and Manager

**6. Thinking from a bank's perspective, which metric should our primary focus be on..**

- ROC AUC
- Precision
- Recall
- F1 Score

F1 Score metric should be our primary focus because both precision and recall is important.

- ROC-AUC : Not good metric to consider as we have highly imbalanced data.
- Precision: Consider when only want to reduce NPA
- Precision: lower precision will lead to losing the business by misppredicting disciplined payees as defaulters
- Recall: lower recall will lead to risk by disbursing the loans to defaulter by mispredicting defaulters as disciplined payees

**7. How does the gap in precision and recall affect the bank?**

- Recall score: 0.96 and Precision score: 0.88. which tells us that there are more false positives than the false negatives.
- If Recall value is low, it means Bank is loosing in opportunity cost.
- If Precision value is low, it means Bank's NPA (defaulters) may increase.

**8. Which were the features that heavily affected the outcome?**

int\_rate, sub\_grade, term, home\_ownership, purpose, application\_type, zipcode (from address), Mortgage Account

**9. Will the results be affected by geographical location? (Yes/No)**

Yes, zipcode has significant impact on the outcome