

# Business Case: LoanTap Logistic Regression

```
In [301]: import pandas as pd
pd.set_option('display.max_columns', 500)
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import re
import warnings
warnings.filterwarnings("ignore")

from statsmodels.stats.outliers_influence import variance_inflation_factor

from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split, KFold, cross_val_score
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report, precision_recall_curve
from sklearn.preprocessing import MinMaxScaler, StandardScaler

from imblearn.over_sampling import SMOTE
```

## Problem Statement

LoanTap is an online platform committed to delivering customized loan products to millennials. They innovate in an otherwise dull loan segment, to deliver instant, flexible loans on consumer friendly terms to salaried professionals and businessmen.

The data science team at LoanTap is building an underwriting layer to determine the creditworthiness of MSMEs as well as individuals.

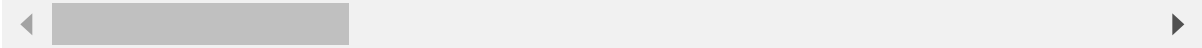
LoanTap different types of loans, This case study will focus on the underwriting process behind Personal Loan only and determine if a credit line should be extended to them. If so, what should the repayment terms be in business recommendations.

Since this is a Classification problem, here both Precision and Recall are important because compromising on precision leads to loss in opportunity to give loans to good customers where we could earn more money and compromising on recall leads to increase in NPA which will affect in profitability of the finance company, so lets build a model we get the best of both.

```
In [2]: df = pd.read_csv('LoanTapData.csv')
df.head()
```

```
Out[2]:
```

	loan_amnt	term	int_rate	installment	grade	sub_grade	emp_title	emp_length	home_
0	10000.0	36 months	11.44	329.48	B	B4	Marketing	10+ years	
1	8000.0	36 months	11.99	265.68	B	B5	Credit analyst	4 years	M
2	15600.0	36 months	10.49	506.97	B	B3	Statistician	< 1 year	
3	7200.0	36 months	6.49	220.65	A	A2	Client Advocate	6 years	
4	24375.0	60 months	17.27	609.33	C	C5	Destiny Management Inc.	9 years	M



```
In [3]: df.shape
```

```
Out[3]: (396030, 27)
```

```
In [4]: df.dtypes
```

```
Out[4]: loan_amnt      float64
term                object
int_rate            float64
installment         float64
grade              object
sub_grade          object
emp_title           object
emp_length          object
home_ownership      object
annual_inc          float64
verification_status object
issue_d             object
loan_status         object
purpose            object
title              object
dti                 float64
earliest_cr_line    object
open_acc            float64
pub_rec             float64
revol_bal           float64
revol_util          float64
total_acc           float64
initial_list_status object
application_type     object
mort_acc            float64
pub_rec_bankruptcies float64
address             object
dtype: object
```

```
In [5]: df['issue_d'] = pd.to_datetime(df['issue_d']) ## converting object types to datetime
df['earliest_cr_line'] = pd.to_datetime(df['earliest_cr_line'])
```

```
In [ ]:
```

```
In [6]: df.describe().T # Statistical summary of numerical columns
```

Out[6]:

	count	mean	std	min	25%	50%	75%
loan_amnt	396030.0	14113.888089	8357.441341	500.00	8000.00	12000.00	20000.00
int_rate	396030.0	13.639400	4.472157	5.32	10.49	13.33	16.00
installment	396030.0	431.849698	250.727790	16.08	250.33	375.43	567.00
annual_inc	396030.0	74203.175798	61637.621158	0.00	45000.00	64000.00	90000.00
dti	396030.0	17.379514	18.019092	0.00	11.28	16.91	22.00
open_acc	396030.0	11.311153	5.137649	0.00	8.00	10.00	14.00
pub_rec	396030.0	0.178191	0.530671	0.00	0.00	0.00	0.00
revol_bal	396030.0	15844.539853	20591.836109	0.00	6025.00	11181.00	19620.00
revol_util	395754.0	53.791749	24.452193	0.00	35.80	54.80	72.00
total_acc	396030.0	25.414744	11.886991	2.00	17.00	24.00	32.00
mort_acc	358235.0	1.813991	2.147930	0.00	0.00	1.00	3.00
pub_rec_bankruptcies	395495.0	0.121648	0.356174	0.00	0.00	0.00	0.00

```
In [7]: df.describe(include='object').T # Statistical summary of categorical columns
```

Out[7]:

	count	unique	top	freq
term	396030	2	36 months	302005
grade	396030	7	B	116018
sub_grade	396030	35	B3	26655
emp_title	373103	173105	Teacher	4389
emp_length	377729	11	10+ years	126041
home_ownership	396030	6	MORTGAGE	198348
verification_status	396030	3	Verified	139563
loan_status	396030	2	Fully Paid	318357
purpose	396030	14	debt_consolidation	234507
title	394275	48817	Debt consolidation	152472
initial_list_status	396030	2	f	238066
application_type	396030	3	INDIVIDUAL	395319
address	396030	393700	USCGC Smith	8

```
In [ ]:
```

In [ ]:

In [8]: *# defaults % increase with rise in interest rates*

```
In [9]: num_cols = df.columns[df.dtypes=='float64']
cat_cols = df.columns[df.dtypes=='O']
date_cols = df.columns[df.dtypes=='datetime64[ns]']
```

In [ ]:

```
In [10]: print("Numeric columns : ", num_cols)
print()
print("Categorical columns : ", cat_cols)
print()
print("datetime columns : ", date_cols)
```

```
Numeric columns : Index(['loan_amnt', 'int_rate', 'installment', 'annual_in
c', 'dti', 'open_acc',
                        'pub_rec', 'revol_bal', 'revol_util', 'total_acc', 'mort_acc',
                        'pub_rec_bankruptcies'],
                        dtype='object')
```

```
Categorical columns : Index(['term', 'grade', 'sub_grade', 'emp_title', 'em
p_length',
                            'home_ownership', 'verification_status', 'loan_status', 'purpose',
                            'title', 'initial_list_status', 'application_type', 'address'],
                            dtype='object')
```

```
datetime columns : Index(['issue_d', 'earliest_cr_line'], dtype='object')
```

In [ ]:

```
In [11]: (100*df.isna().sum()/df.shape[0]).round(2)    # Percentage of nulls
```

```
Out[11]: loan_amnt      0.00
term      0.00
int_rate  0.00
installment  0.00
grade     0.00
sub_grade 0.00
emp_title  5.79
emp_length 4.62
home_ownership 0.00
annual_inc 0.00
verification_status 0.00
issue_d    0.00
loan_status 0.00
purpose    0.00
title      0.44
dti        0.00
earliest_cr_line 0.00
open_acc   0.00
pub_rec    0.00
revol_bal  0.00
revol_util 0.07
total_acc  0.00
initial_list_status 0.00
application_type 0.00
mort_acc   9.54
pub_rec_bankruptcies 0.14
address     0.00
dtype: float64
```


```
In [12]: def parse_numeric(x):
         return '' if pd.isna(x) else re.findall(r'\d+',x)[0]
```

```
In [13]: df['term'].unique()
```

```
Out[13]: array([' 36 months', ' 60 months'], dtype=object)
```

```
In [14]: sns.countplot(df['term'])
```

```
Out[14]: <AxesSubplot:xlabel='term', ylabel='count'>
```



```
In [15]: df['term'] = df['term'].apply(lambda x : parse_numeric(x)).apply(int) # con
```


```
In [ ]:
```

```
In [16]: df['emp_length'].unique() # We can parse to remove text from the column and l
```

```
Out[16]: array(['10+ years', '4 years', '< 1 year', '6 years', '9 years',  
                '2 years', '3 years', '8 years', '7 years', '5 years', '1 year',  
                nan], dtype=object)
```

```
In [17]: plt.figure(figsize=(9,4))
sns.countplot(df['emp_length'],hue=df['loan_status'])
```

```
Out[17]: <AxesSubplot:xlabel='emp_length', ylabel='count'>
```



```
In [18]: df['emp_length'] = pd.to_numeric(df['emp_length'].apply(lambda x : parse_num
```

```
In [19]: df['emp_length'].unique()
```

```
Out[19]: array([10.,  4.,  1.,  6.,  9.,  2.,  3.,  8.,  7.,  5., nan])
```

```
In [ ]:
```

**Filling missing values of emp\_len with difference between initial credit line year to issue year as this could give us the approximate employment years**

```
In [20]: temp_emp_length_fill = df['issue_d'].dt.year-df['earliest_cr_line'].dt.year
```

```
In [21]: temp_emp_length_fill[temp_emp_length_fill>10]=10
```



```
In [22]: temp_emp_length_fill.value_counts()
```

```
Out[22]: 10    329781
          9     16382
          8     14369
          7     12281
          6      9507
          5      6969
          4      4997
          3      1744
          dtype: int64
```

```
In [23]: df['emp_length'] = df['emp_length'].fillna(temp_emp_length_fill)
```

```
In [24]: df['grade'].unique() # As grade can be formulated to label encoding we label
```

```
Out[24]: array(['B', 'A', 'C', 'E', 'D', 'F', 'G'], dtype=object)
```

```
In [25]: grade_dict = dict(zip(list('ABCDEFG'),[i for i in range(7)]))
df['grade'] = df['grade'].replace(grade_dict)
```

```
In [ ]:
```

```
In [26]: df['sub_grade'].unique()
```

```
Out[26]: array(['B4', 'B5', 'B3', 'A2', 'C5', 'C3', 'A1', 'B2', 'C1', 'A5', 'E4',
               'A4', 'A3', 'D1', 'C2', 'B1', 'D3', 'D5', 'D2', 'E1', 'E2', 'E5',
               'F4', 'E3', 'D4', 'G1', 'F5', 'G2', 'C4', 'F1', 'F3', 'G5', 'G4',
               'F2', 'G3'], dtype=object)
```

```
In [308]: df['sub_grade'] = df['sub_grade'].str[1].apply(int) ## subgrade is extension of
```


```
In [ ]:
```

```
In [28]: df['home_ownership'].unique()
```

```
Out[28]: array(['RENT', 'MORTGAGE', 'OWN', 'OTHER', 'NONE', 'ANY'], dtype=object)
```

```
In [29]: sns.countplot(df['home_ownership'])
```

```
Out[29]: <AxesSubplot:xlabel='home_ownership', ylabel='count'>
```



```
In [30]: home_ownership_dict = {'NONE':'OTHER','ANY':'OTHER'} # we will replace None and ANY with OTHER
df['home_ownership'] = df['home_ownership'].replace(home_ownership_dict)
```


```
In [ ]:
```

```
In [31]: df['verification_status'].unique()
```

```
Out[31]: array(['Not Verified', 'Source Verified', 'Verified'], dtype=object)
```

```
In [32]: sns.countplot(df['verification_status'], hue=df['loan_status'])
```

```
Out[32]: <AxesSubplot:xlabel='verification_status', ylabel='count'>
```



```
In [33]: verification_status_dict = dict(zip(['Not Verified', 'Source Verified', 'Verified'], df['verification_status'].replace(verification_status_dict)))
```


```
In [ ]:
```

```
In [34]: df['loan_status'].unique()
```

```
Out[34]: array(['Fully Paid', 'Charged Off'], dtype=object)
```

```
In [35]: sns.countplot(df['loan_status'])
```

```
Out[35]: <AxesSubplot:xlabel='loan_status', ylabel='count'>
```



```
In [36]: loan_status_dict = dict(zip(['Fully Paid', 'Charged Off'],[0,1]))  
df['loan_status'] = df['loan_status'].replace(loan_status_dict)
```

```
In [ ]:
```

```
In [37]: df['initial_list_status'].value_counts()
```

```
Out[37]: f    238066  
        w    157964  
        Name: initial_list_status, dtype: int64
```

```
In [38]: initial_list_status_dict = dict(zip(['w', 'f'],[1,0]))  
df['initial_list_status'] = df['initial_list_status'].replace(initial_list_status_dict)
```

```
In [ ]:
```

```
In [39]: df['address'].unique()
```

```
Out[39]: array(['0174 Michelle Gateway\r\nMendozaberg, OK 22690',
               '1076 Carney Fort Apt. 347\r\nLoganmouth, SD 05113',
               '87025 Mark Dale Apt. 269\r\nNew Sabrina, WV 05113', ...,
               '953 Matthew Points Suite 414\r\nReedfort, NY 70466',
               '7843 Blake Freeway Apt. 229\r\nNew Michael, FL 29597',
               '787 Michelle Causeway\r\nBriannaton, AR 48052'], dtype=object)
```

```
In [40]: df['address'].value_counts()
```

```
Out[40]: USCGC Smith\r\nFPO AE 70466      8
USS Johnson\r\nFPO AE 48052      8
USNS Johnson\r\nFPO AE 05113      8
USS Smith\r\nFPO AP 70466      8
USNS Johnson\r\nFPO AP 48052      7
..
455 Tricia Cove\r\nAustinbury, FL 00813      1
7776 Flores Fall\r\nFernandezshire, UT 05113      1
6577 Mia Harbors Apt. 171\r\nRobertshire, OK 22690      1
8141 Cox Greens Suite 186\r\nMadisonstad, VT 05113      1
787 Michelle Causeway\r\nBriannaton, AR 48052      1
Name: address, Length: 393700, dtype: int64
```

```
In [41]: df['address'] = df['address'].str[-5:]
```

```
In [42]: df['address'].value_counts()
```

```
Out[42]: 70466      56985
30723      56546
22690      56527
48052      55917
00813      45824
29597      45471
05113      45402
11650      11226
93700      11151
86630      10981
Name: address, dtype: int64
```

```
In [ ]:
```

```
In [167]: df['purpose'].unique()
```

```
Out[167]: array(['vacation', 'debt_consolidation', 'credit_card',
               'home_improvement', 'small_business', 'major_purchase', 'other',
               'medical', 'wedding', 'car', 'moving', 'house', 'educational',
               'renewable_energy'], dtype=object)
```

```
In [159]: (100*df.isna().sum()/df.shape[0]).round(2)
```

```
Out[159]: loan_amnt      0.00
term      0.00
int_rate  0.00
installment 0.00
grade     0.00
sub_grade 0.00
emp_title  0.00
emp_length 0.00
home_ownership 0.00
annual_inc 0.00
verification_status 0.00
issue_d    0.00
loan_status 0.00
purpose    0.00
title      0.44
dti        0.00
earliest_cr_line 0.00
open_acc   0.00
pub_rec    0.00
revol_bal  0.00
revol_util 0.00
total_acc  0.00
initial_list_status 0.00
application_type 0.00
mort_acc   0.00
pub_rec_bankruptcies 0.00
address    0.00
revol_util_na 0.00
mort_acc_na 0.00
dtype: float64
```

```
In [44]: df['revol_util_na']=df['revol_util'].isna().apply(int)
```

```
In [45]: df['revol_util'] = df['revol_util'].fillna(0)
```

```
In [ ]:
```

```
In [46]: df['emp_title'] = df['emp_title'].str.lower().str.strip().fillna('other')
```


```
In [ ]:
```

```
In [54]: df['pub_rec_bankruptcies'].value_counts()
```

```
Out[54]: 0.0    350380
         1.0    42790
         2.0    1847
         3.0     351
         4.0     82
         5.0     32
         6.0      7
         7.0      4
         8.0      2
         Name: pub_rec_bankruptcies, dtype: int64
```

```
In [49]: pd.crosstab(columns = df["loan_status"],
                    index=df['pub_rec_bankruptcies'],
                    normalize="index").plot(kind="bar")
```

```
Out[49]: <AxesSubplot:xlabel='pub_rec_bankruptcies'>
```




```
In [62]: df['pub_rec'].value_counts()
```

```
Out[62]: 0.0    338272
         1.0    49739
         2.0     5476
         3.0    1521
         4.0     527
         5.0     237
         6.0     122
         7.0      56
         8.0      34
         9.0      12
        10.0       11
        11.0        8
        13.0        4
        12.0        4
        19.0        2
        40.0        1
        17.0        1
        86.0        1
        24.0        1
        15.0        1
        Name: pub_rec, dtype: int64
```

```
In [53]: sns.countplot(df['pub_rec'])
```

```
Out[53]: <AxesSubplot:xlabel='pub_rec', ylabel='count'>
```





```
In [67]: df['pub_rec_bankruptcies'] = df['pub_rec_bankruptcies'].fillna(df['pub_rec'])
```

```
In [ ]:
```

```
In [ ]:
```

```
In [70]: df['mort_acc'].value_counts()
```

```
Out[70]: 0.0      139777
         1.0       60416
         2.0       49948
         3.0       38049
         4.0       27887
         5.0       18194
         6.0       11069
         7.0        6052
         8.0        3121
         9.0       1656
        10.0        865
        11.0        479
        12.0        264
        13.0        146
        14.0        107
        15.0         61
        16.0         37
        17.0         22
        18.0         18
        19.0          5
```

```
In [120]: df['mort_acc_na'] = df['mort_acc'].isna().apply(int)
```

```
In [122]: df['home_ownership'].unique()
```

```
Out[122]: array(['RENT', 'MORTGAGE', 'OWN', 'OTHER'], dtype=object)
```

```
In [124]: df['purpose'].unique()
```


```
Out[124]: array(['vacation', 'debt_consolidation', 'credit_card',
                 'home_improvement', 'small_business', 'major_purchase', 'other',
                 'medical', 'wedding', 'car', 'moving', 'house', 'educational',
                 'renewable_energy'], dtype=object)
```

```
In [155]: def mort_acc_fillna(row): # filling the mortgage with the best possible value
    if pd.isnull(row['mort_acc']):
        if row['home_ownership'] in ['MORTGAGE', 'OTHER', 'OWN'] or row['purpose'] in ['MORTGAGE']:
            return 1
        else:
            return 0
    else:
        return row['mort_acc']
```

```
In [158]: df['mort_acc'] = df.apply(mort_acc_fillna, axis=1)
```


```
In [71]: pd.crosstab(columns = df["loan_status"],
                    index=df['mort_acc'],
                    normalize="index").plot(kind="bar")
```

```
Out[71]: <AxesSubplot:xlabel='mort_acc'>
```



```
In [72]: pd.crosstab(columns = df["loan_status"],
                    index=df['home_ownership'],
                    normalize="index").plot(kind="bar")
```

Out[72]: <AxesSubplot:xlabel='home\_ownership'>



In [ ]:

```
In [84]: df[df['home_ownership']=='OWN'][['mort_acc', 'total_acc', 'open_acc']].corr()
```

Out[84]:

	mort_acc	total_acc	open_acc
mort_acc	1.000000	0.352143	0.104398
total_acc	0.352143	1.000000	0.696238
open_acc	0.104398	0.696238	1.000000

```
In [82]: (df['open_acc']>=df['mort_acc']).value_counts()
```

Out[82]: True 355208  
False 40822  
dtype: int64


In [ ]:

In [ ]:

In [ ]:


```
In [173]: pd.crosstab(columns = df["loan_status"],  
                    index=pd.qcut(df["annual_inc"],15),  
                    normalize="index").plot(kind="bar")
```

Out[173]: <AxesSubplot:xlabel='annual\_inc'>




```
In [164]: pd.crosstab(columns = df["loan_status"],
                    index=pd.qcut(df["int_rate"],15),
                    normalize="index").plot(kind="bar")
```

Out[164]: <AxesSubplot:xlabel='int\_rate'>



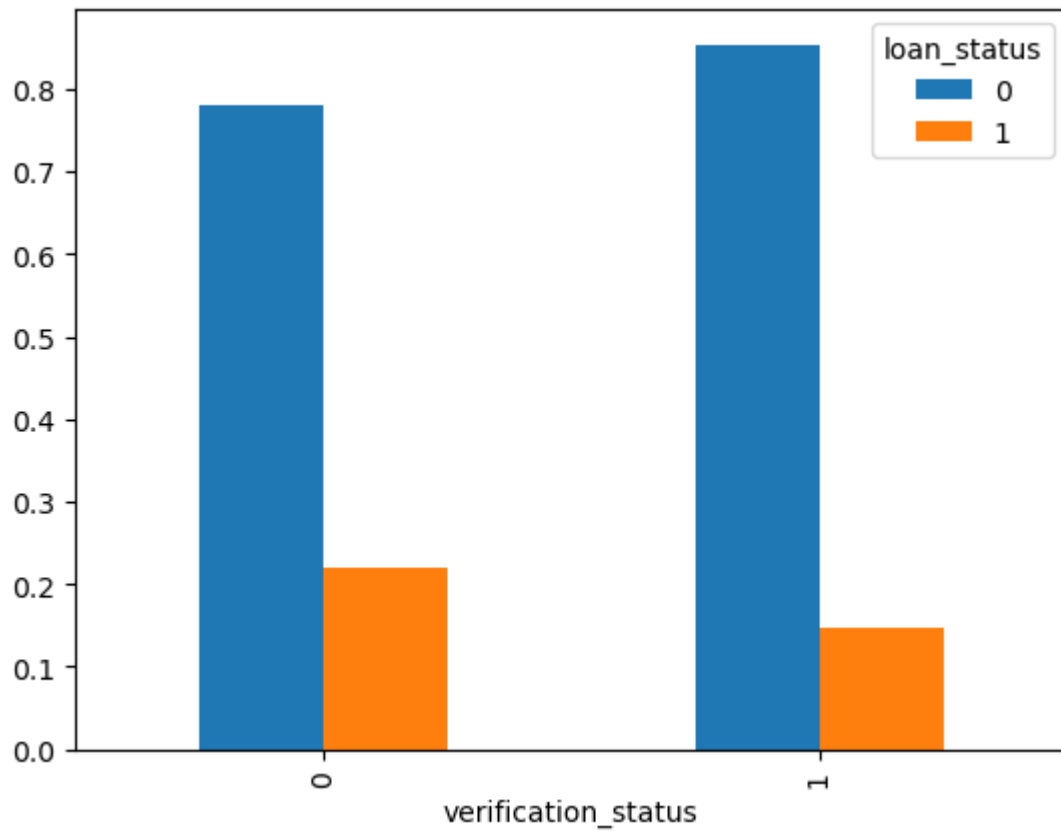
```
In [165]: pd.crosstab(columns = df["loan_status"],  
                    index=df["grade"],  
                    normalize="index").plot(kind="bar")
```

Out[165]: <AxesSubplot:xlabel='grade'>



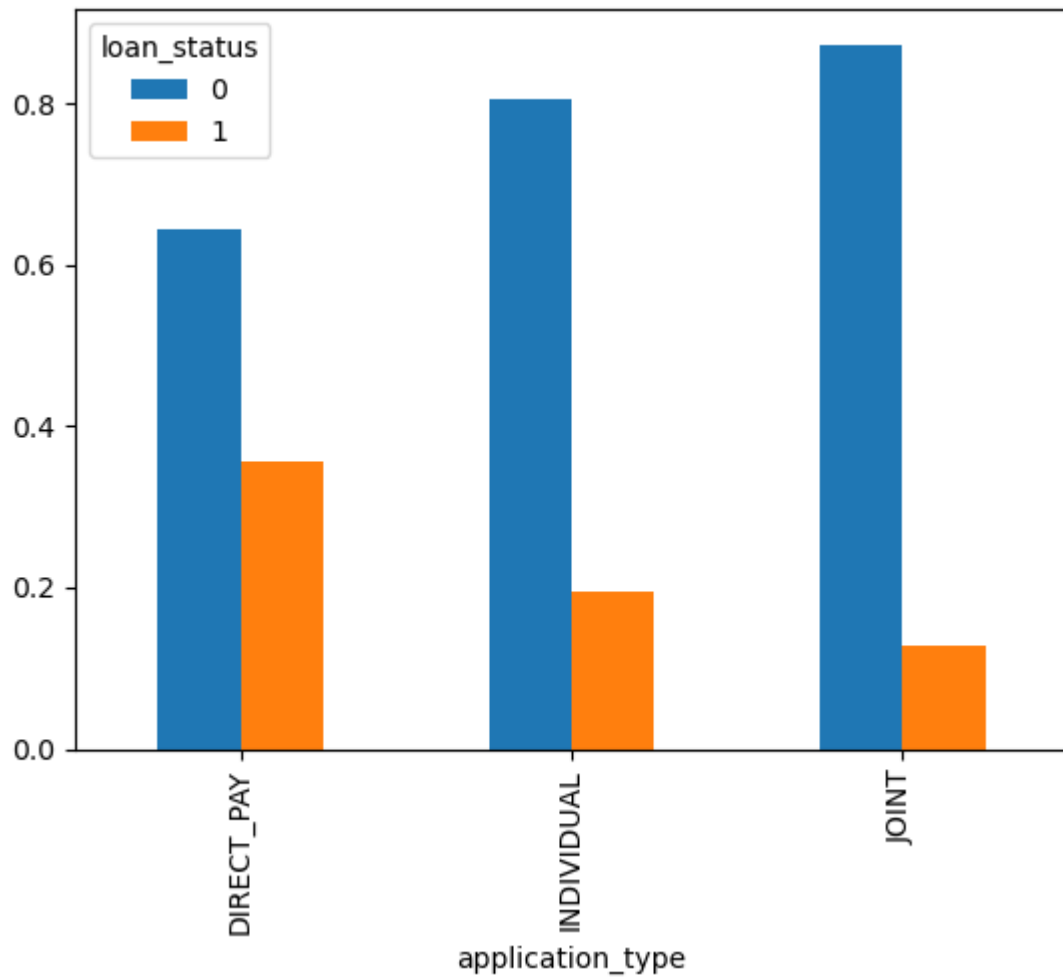
```
In [166]: pd.crosstab(columns = df["loan_status"],  
                    index=df["verification_status"],  
                    normalize="index").plot(kind="bar")
```

Out[166]: <AxesSubplot:xlabel='verification\_status'>



```
In [167]: pd.crosstab(columns = df["loan_status"],  
                    index=df["application_type"],  
                    normalize="index").plot(kind="bar")
```


Out[167]: <AxesSubplot:xlabel='application\_type'>






```
In [172]: pd.crosstab(columns = df["loan_status"],
                    index=df['mort_acc'],
                    normalize="index").plot(kind="bar")
```

Out[172]: <AxesSubplot:xlabel='mort\_acc'>



```
In [175]: pd.crosstab(columns = df["loan_status"],  
                    index=df["application_type"],  
                    normalize="index").plot(kind="bar")
```

Out[175]: <AxesSubplot:xlabel='application\_type'>



In [ ]:

In [ ]:

```
In [92]: mort_acc_na = df['mort_acc'].apply(mort)
```

```
In [89]: pd.crosstab(columns = df[df['loan_status']=='Charged Off']['home_ownership'],
                    index=df[df['loan_status']=='Charged Off']['mort'],
                    normalize="index")
```

```
Out[89]:
```

home_ownership	MORTGAGE	NONE	OTHER	OWN	RENT
mort					
-1	0.406126	0.000000	0.001622	0.082703	0.509550
0	0.095063	0.000060	0.000121	0.125554	0.779202
1	0.724375	0.000128	0.000077	0.081712	0.193707

```
In [94]: pd.crosstab(columns = df["home_ownership"],
                    index=df['mort'],
                    normalize="index")
```

```
Out[94]:
```

home_ownership	ANY	MORTGAGE	NONE	OTHER	OWN	RENT
mort						
-1	0.000000	0.440323	0.000053	0.002064	0.078502	0.479058
0	0.000000	0.104209	0.000064	0.000107	0.122896	0.772724
1	0.000014	0.765090	0.000092	0.000087	0.080569	0.154149

```
In [93]: pd.crosstab(columns = df["home_ownership"],
                    index=mort_acc_na)
```


```
Out[93]:
```

home_ownership	MORTGAGE	OTHER	OWN	RENT
mort_acc				
-1	16642	80	2967	18106
0	14566	24	17178	108009
1	167140	42	17601	33675

```
In [104]: df_ = df[df["home_ownership"]=='OWN']
```

```
In [108]: pd.crosstab(columns = df["loan_status"],
                    index=(df["emp_length"]),
                    normalize="index").plot(kind="bar")
```

Out[108]: <AxesSubplot:xlabel='emp\_length'>



```
In [94]: pd.crosstab(columns = df["home_ownership"],
                    index=mort_acc_na,
                    normalize="columns")
```

Out[94]:

home_ownership	MORTGAGE	OTHER	OWN	RENT
mort_acc				
-1	0.083903	0.547945	0.078604	0.113311
0	0.073437	0.164384	0.455095	0.675943
1	0.842660	0.287671	0.466301	0.210745

```
In [93]: df[df['mort_acc'].isna()].groupby(['home_ownership'])['loan_status'].count()
```


```
Out[93]: home_ownership
MORTGAGE    16642
NONE         2
OTHER        78
OWN         2967
RENT       18106
Name: loan_status, dtype: int64
```

```
In [ ]: pd.crosstab(columns = df["loan_status"],
                    index=df['emp_length_na'].isna(),
                    normalize="index").plot(kind="bar",grid=True)
```

```
In [102]: df['emp_length_na'] = df['emp_length'].isna()
```

```
In [101]: pd.crosstab(columns = df["loan_status"],
                    index=df['emp_length'].isna(),
                    normalize="index").plot(kind="bar",grid=True)
```

```
Out[101]: <AxesSubplot:xlabel='emp_length'>
```



```
In [110]: df['purpose']=='house'
```

```
Out[110]: array(['vacation', 'debt_consolidation', 'credit_card',  
                'home_improvement', 'small_business', 'major_purchase', 'other',  
                'medical', 'wedding', 'car', 'moving', 'house', 'educational',  
                'renewable_energy'], dtype=object)
```

```
In [116]: df['purpose'].value_counts(normalize=True)*100
```

```
Out[116]: debt_consolidation    59.214453  
credit_card                    20.962806  
home_improvement              6.067722  
other                         5.349342  
major_purchase                2.219529  
small_business                1.439537  
car                           1.186021  
medical                       1.059516  
moving                        0.720652  
vacation                      0.619145  
house                         0.555766  
wedding                       0.457541  
renewable_energy              0.083075  
educational                   0.064894  
Name: purpose, dtype: float64
```

```
In [117]: df[df['mort_acc'].isna()][ 'purpose'].value_counts(normalize=True)*100
```

```
Out[117]: debt_consolidation    47.712660  
credit_card                    14.075936  
other                         9.564757  
home_improvement              7.082947  
major_purchase                5.125017  
small_business                4.651409  
car                           3.728006  
wedding                       2.262204  
medical                       1.656304  
moving                        1.349385  
house                         1.010716  
vacation                      0.870486  
educational                   0.677338  
renewable_energy              0.232835  
Name: purpose, dtype: float64
```

```
In [113]: (df['mort_acc'].isna()).value_counts()
```

```
Out[113]: False    358235  
          True      37795  
          Name: mort_acc, dtype: int64
```

In [ ]:

In [ ]:

In [176]: df.head()

Out[176]:

	loan_amnt	term	int_rate	installment	grade	sub_grade	emp_title	emp_length	home_ov
0	10000.0	36	11.44	329.48	1	4	marketing	10.0	
1	8000.0	36	11.99	265.68	1	5	credit analyst	4.0	MOI
2	15600.0	36	10.49	506.97	1	3	statistician	1.0	
3	7200.0	36	6.49	220.65	0	2	client advocate	6.0	
4	24375.0	60	17.27	609.33	2	5	destiny management inc.	9.0	MOI

In [248]: for col in ['annual\_inc', 'dti', 'open\_acc', 'revol\_bal', 'revol\_util']:  
sns.boxplot(df[col])  
plt.show()



In [ ]:

```
In [192]: df['annual_inc'].shape[0]-df.loc[df['annual_inc']<260000,'annual_inc'].shape[0]
```


```
Out[192]: 3194
```

```
In [210]: df.loc[df['annual_inc']<1100000,'loan_status'].value_counts()
```

```
Out[210]: 0    50  
         1     7  
         Name: loan_status, dtype: int64
```

```
In [209]: sns.histplot((df.loc[df['annual_inc']<110000,'annual_inc']),bins=15)
```


```
Out[209]: <AxesSubplot:xlabel='annual_inc', ylabel='Count'>
```





```
In [177]: sns.histplot(df['loan_amnt'],bins=15)
```

```
Out[177]: <AxesSubplot:xlabel='loan_amnt', ylabel='Count'>
```



```
In [262]: df_1=df.copy(deep=True)
```

```
In [263]: df_1.to_csv("preprocessed-data.csv",index=False)
```

```
In [250]: df.shape
```

```
Out[250]: (396030, 29)
```

```
In [251]: data = df[(df['annual_inc']<110000) & (df['dti']<60) & (df['open_acc']<60) &
```

```
In [252]: data.shape
```

```
Out[252]: (395945, 29)
```

```
In [254]: data['revol_bal'] = np.log1p(data['revol_bal'])
```

In [ ]:

In [261]: data.head()

Out[261]:

	loan_amnt	term	int_rate	installment	grade	sub_grade	emp_length	home_ownership	ann
0	10000.0	36	11.44	329.48	1	4	10.0	RENT	1
1	8000.0	36	11.99	265.68	1	5	4.0	MORTGAGE	6
2	15600.0	36	10.49	506.97	1	3	1.0	RENT	4
3	7200.0	36	6.49	220.65	0	2	6.0	RENT	5
4	24375.0	60	17.27	609.33	2	5	9.0	MORTGAGE	5

In [257]: data.drop(columns=['emp\_title', 'title'], inplace=True)

In [260]: data['issue\_d'] = data['issue\_d'].dt.year  
data['earliest\_cr\_line'] = data['earliest\_cr\_line'].dt.year

In [264]: dummies = ['home\_ownership', 'purpose', 'application\_type', 'address']  
data = pd.get\_dummies(data, columns=dummies, drop\_first=True)

In [265]: data.shape

Out[265]: (395945, 50)

In [266]: data.head()

Out[266]:


	loan_amnt	term	int_rate	installment	grade	sub_grade	emp_length	annual_inc	verification
0	10000.0	36	11.44	329.48	1	4	10.0	117000.0	
1	8000.0	36	11.99	265.68	1	5	4.0	65000.0	
2	15600.0	36	10.49	506.97	1	3	1.0	43057.0	
3	7200.0	36	6.49	220.65	0	2	6.0	54000.0	
4	24375.0	60	17.27	609.33	2	5	9.0	55000.0	

In [267]: data['pub\_rec'].apply(lambda x :0 if x==0 else 1).value\_counts()

Out[267]: 0 338194  
1 57751  
Name: pub\_rec, dtype: int64



```
In [275]: plt.figure(figsize=(30, 25))
sns.heatmap(pearson_corr, annot=True, cmap='viridis')
plt.show()
```



```
In [276]: data.describe()
```

```
Out[276]:
```

	loan_amnt	term	int_rate	installment	grade	sub_gra
count	395945.000000	395945.000000	395945.000000	395945.000000	395945.000000	395945.000000
mean	14112.684969	41.698185	13.639617	471.811657	1.822389	2.9717
std	8356.309080	10.212120	4.472038	250.683221	1.333799	1.4067
min	500.000000	36.000000	5.320000	16.080000	0.000000	1.0000
25%	8000.000000	36.000000	10.490000	250.330000	1.000000	2.0000
50%	12000.000000	36.000000	13.330000	375.430000	2.000000	3.0000
75%	20000.000000	36.000000	16.490000	567.300000	3.000000	4.0000
max	40000.000000	60.000000	30.990000	1533.810000	6.000000	5.0000

```
In [277]: data.isna().sum()
```

```
Out[277]: loan_amnt      0
term      0
int_rate  0
installment  0
grade     0
sub_grade  0
emp_length  0
annual_inc  0
verification_status  0
issue_d     0
loan_status  0
dti         0
earliest_cr_line  0
open_acc    0
pub_rec     0
revol_bal   0
revol_util  0
total_acc   0
initial_list_status  0
^
```

```
In [ ]:
```

```
In [278]: X = data.drop(columns=['loan_status'])
y = data['loan_status']
```

```
In [279]: X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.25, stratify=
```

```
In [280]: X_train.shape,X_test.shape,y_train.shape,y_test.shape
```

```
Out[280]: ((296958, 52), (98987, 52), (296958,), (98987,))
```

```
In [281]: scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

```
In [283]: model = LogisticRegression(max_iter=1000)
model.fit(X_train,y_train)
y_pred = model.predict(X_test)
print("Train score : " , model.score(X_train,y_train))
print("Test score : " , model.score(X_test,y_test))
```

```
Train score : 0.8883680520477644
Test score : 0.8910968106923132
```

```
In [288]: model = LogisticRegression(max_iter=1000,class_weight='balanced')
model.fit(X_train,y_train)
y_pred = model.predict(X_test)
print("Train score : " , model.score(X_train,y_train))
print("Test score : " , model.score(X_test,y_test))
```

```
Train score : 0.8100707844206925
Test score : 0.809783102831685
```

In [ ]:

In [ ]:

```
In [285]: print(confusion_matrix(y_test, y_test))
```

```
[[79572    0]
 [    0 19415]]
```

```
In [286]: print(confusion_matrix(y_test, y_pred))
```

```
[[78779    793]
 [ 9987   9428]]
```

```
In [289]: print(confusion_matrix(y_test, y_pred))
```

```
[[64751 14821]
 [ 4008 15407]]
```

In [ ]:

In [ ]:

```
In [287]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.89	0.99	0.94	79572
1	0.92	0.49	0.64	19415
accuracy			0.89	98987
macro avg	0.90	0.74	0.79	98987
weighted avg	0.89	0.89	0.88	98987

```
In [290]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.94	0.81	0.87	79572
1	0.51	0.79	0.62	19415
accuracy			0.81	98987
macro avg	0.73	0.80	0.75	98987
weighted avg	0.86	0.81	0.82	98987


```
In [291]: def precision_recall_curve_plot(y_test, pred_proba_c1):
    precisions, recalls, thresholds = precision_recall_curve(y_test, pred_proba_c1)

    threshold_boundary = thresholds.shape[0]
    # plot precision
    plt.plot(thresholds, precisions[0:threshold_boundary], linestyle='--', label='precision')
    # plot recall
    plt.plot(thresholds, recalls[0:threshold_boundary], label='recalls')

    start, end = plt.xlim()
    plt.xticks(np.round(np.arange(start, end, 0.1), 2))

    plt.xlabel('Threshold Value'); plt.ylabel('Precision and Recall Value')
    plt.legend(); plt.grid()
    plt.show()

precision_recall_curve_plot(y_test, model.predict_proba(X_test)[:,-1])
```



In [ ]:






```
In [298]: lr1 = LogisticRegression(max_iter=5000)
lr1.fit(X_train_res, y_train_res)
predictions = lr1.predict(X_test)

# Classification Report
print(classification_report(y_test, predictions))
```

	precision	recall	f1-score	support
0	0.94	0.81	0.87	79572
1	0.51	0.79	0.62	19415
accuracy			0.81	98987
macro avg	0.72	0.80	0.75	98987
weighted avg	0.86	0.81	0.82	98987

```
In [303]: precision_recall_curve_plot(y_test, lr1.predict_proba(X_test)[:,-1])
```



In [ ]:

In [ ]:

```
In [304]: logreg = LogisticRegression(max_iter=2000,class_weight='balanced')
```

```
In [305]: X = scaler.fit_transform(X_train)

kfold = KFold(n_splits=5)
accuracy = np.mean(cross_val_score(logreg, X_train, y_train, cv=kfold, scoring='accuracy'))
print("Cross Validation accuracy: {:.3f}".format(accuracy))
```

Cross Validation accuracy: 0.810

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```